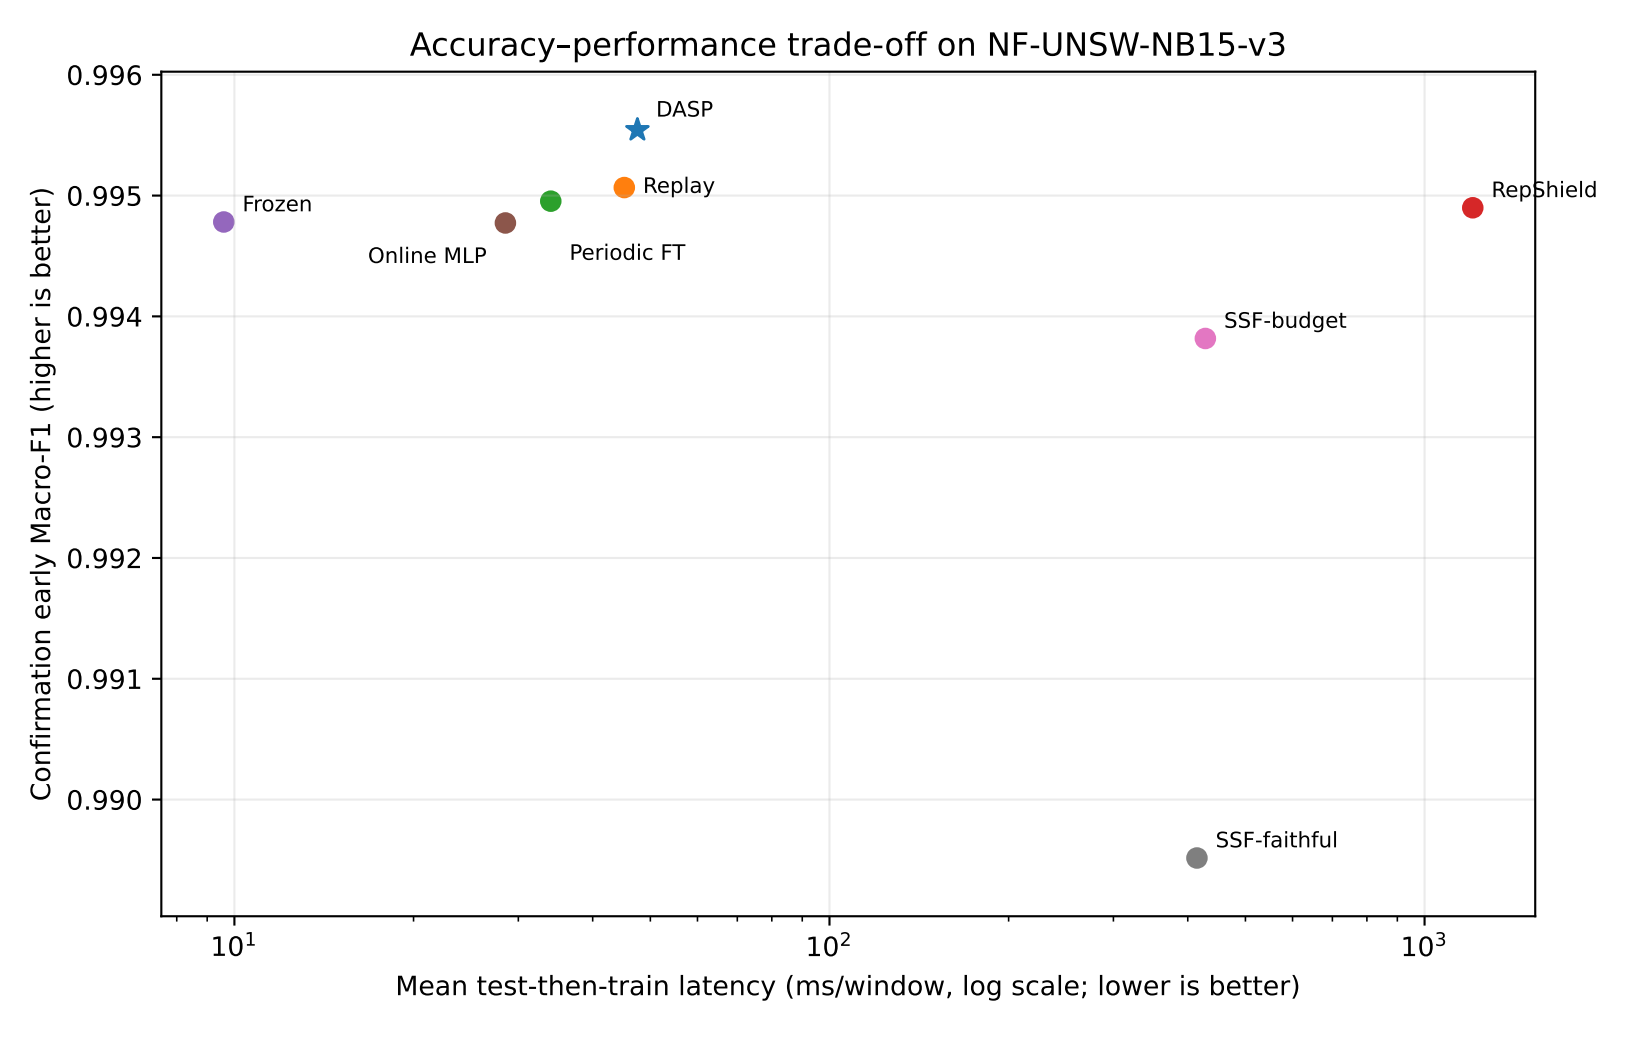
Accuracy–performance trade-off on NF-UNSW-NB15-v3
Confirmation early Macro-F1 (higher is better)
Mean test-then-train latency (ms/window, log scale; lower is better)
0.996
0.995
0.994
0.993
0.992
0.991
0.990
10^1
10^2
10^3
DASP
Replay
Frozen
Online MLP
Periodic FT
RepShield
SSF-budget
SSF-faithful